

TEXT ANALYSIS: MATCHING AND LINKING FOR JOINING DATA

Austin Alleman

November 20, 2019

WORKSHOP RESOURCES

All materials, including these slides and R scripts are available here:

https://github.com/allemanau/NUIT_text_matching_workshop

allemanau / **NUIT_text_matching_workshop**

Code Issues 0 Pull requests 0 Actions Projects 0 Wiki Security Insights Settings

Materials for the RCS text matching workshop.

Manage topics

5 commits 1 branch 0 packages 0 releases 1 contributor

Branch: master New pull request Create new file Upload files Find file Clone or download

Austin Alleman and Austin Alleman PPTX added. Latest commit 4212073 4 hours ago

.DS_Store PPTX added. 4 hours ago

.gitignore Incremental commit. yesterday

INTRODUCTIONS

Who am I?

- Data science research consultant with NUIT RCS
- Fuzzy logic evangelist and practitioner

Who are you?

- Introduce yourself to your neighbors!
- What do you work on, and what brings you here?

WORKSHOP GOALS

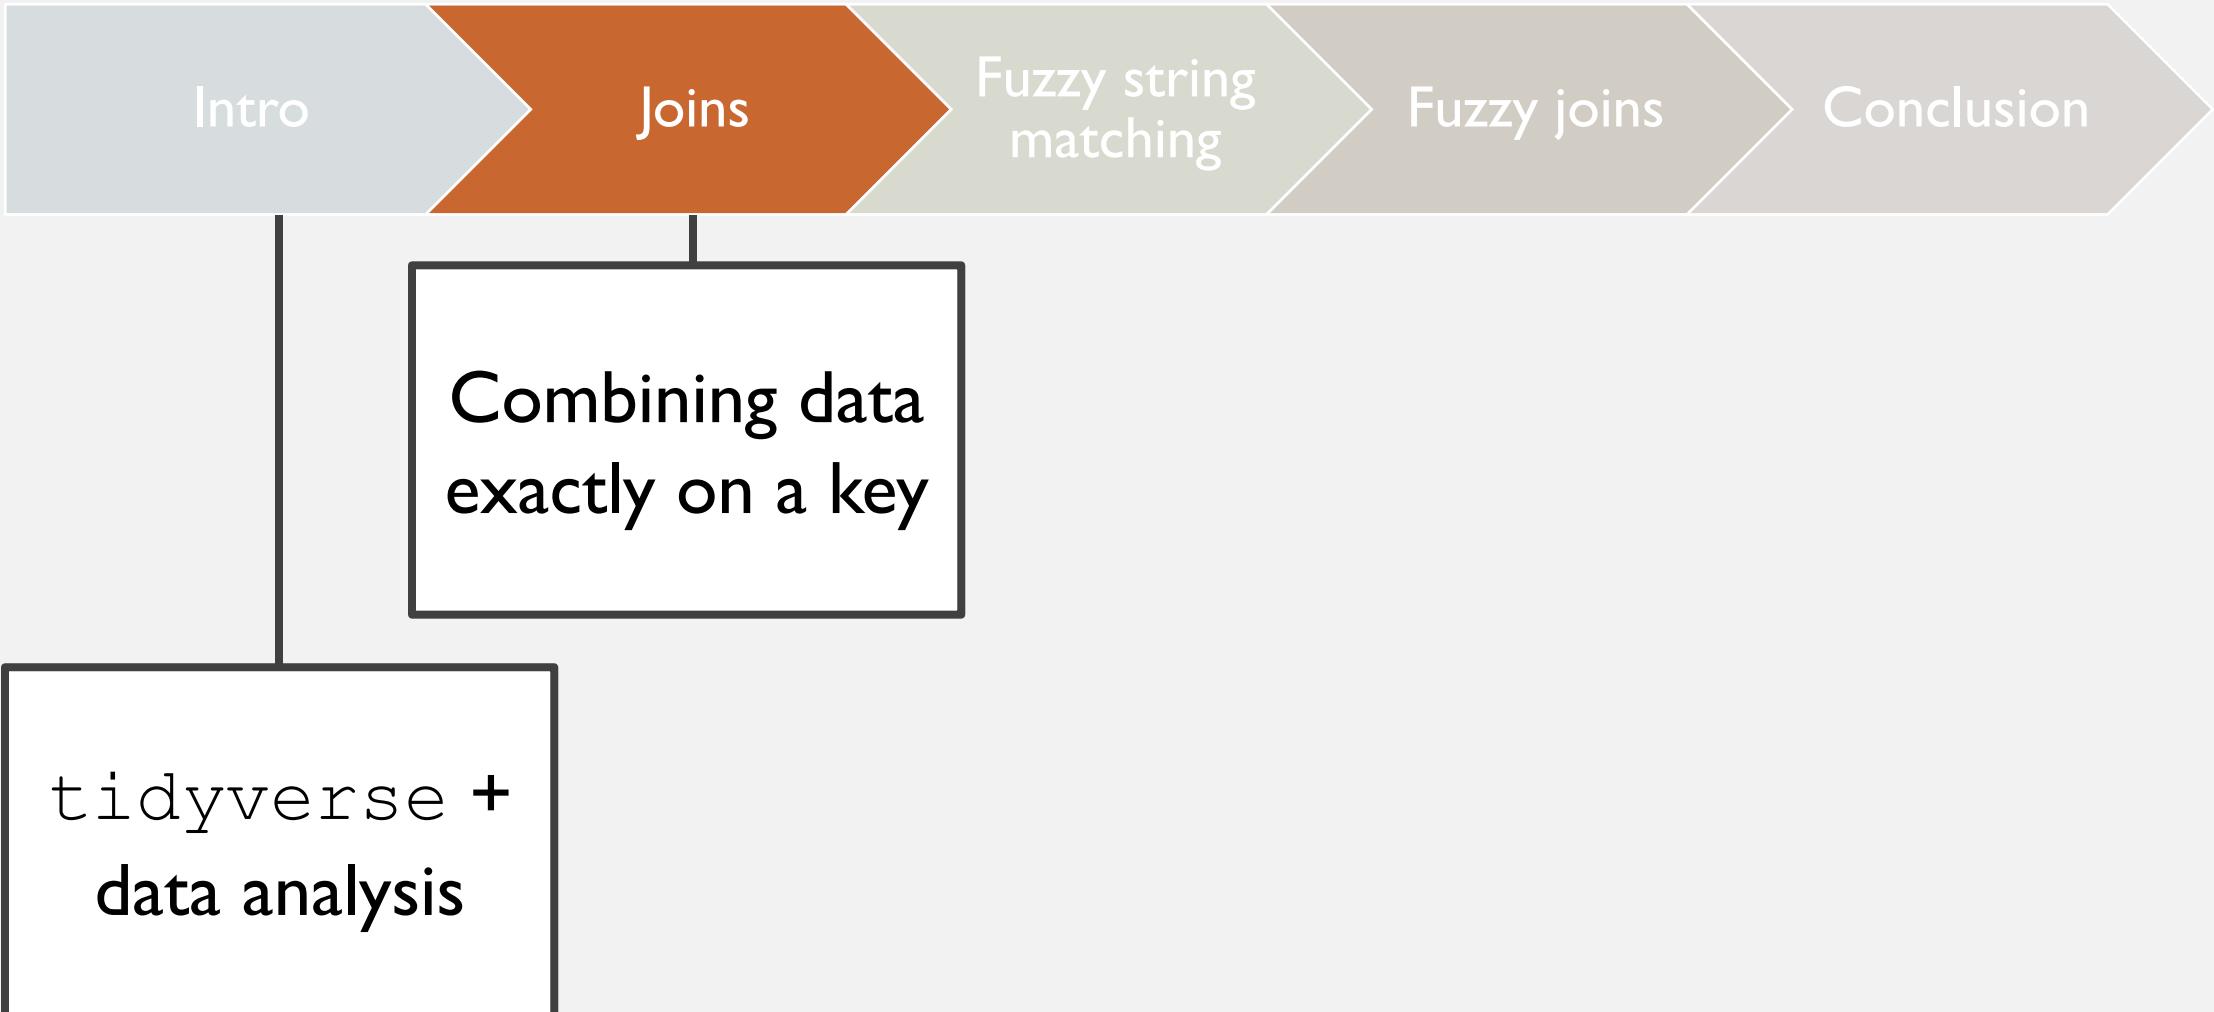
- review joins for combining data sets and highlight cases where exact matching will fail
- discuss some of the mechanics underlying fuzzy matching techniques and parameter tuning
- use R's tidyverse, stringdist, and fuzzyjoin libraries to work through an example on real-world data

ORDER OF OPERATIONS

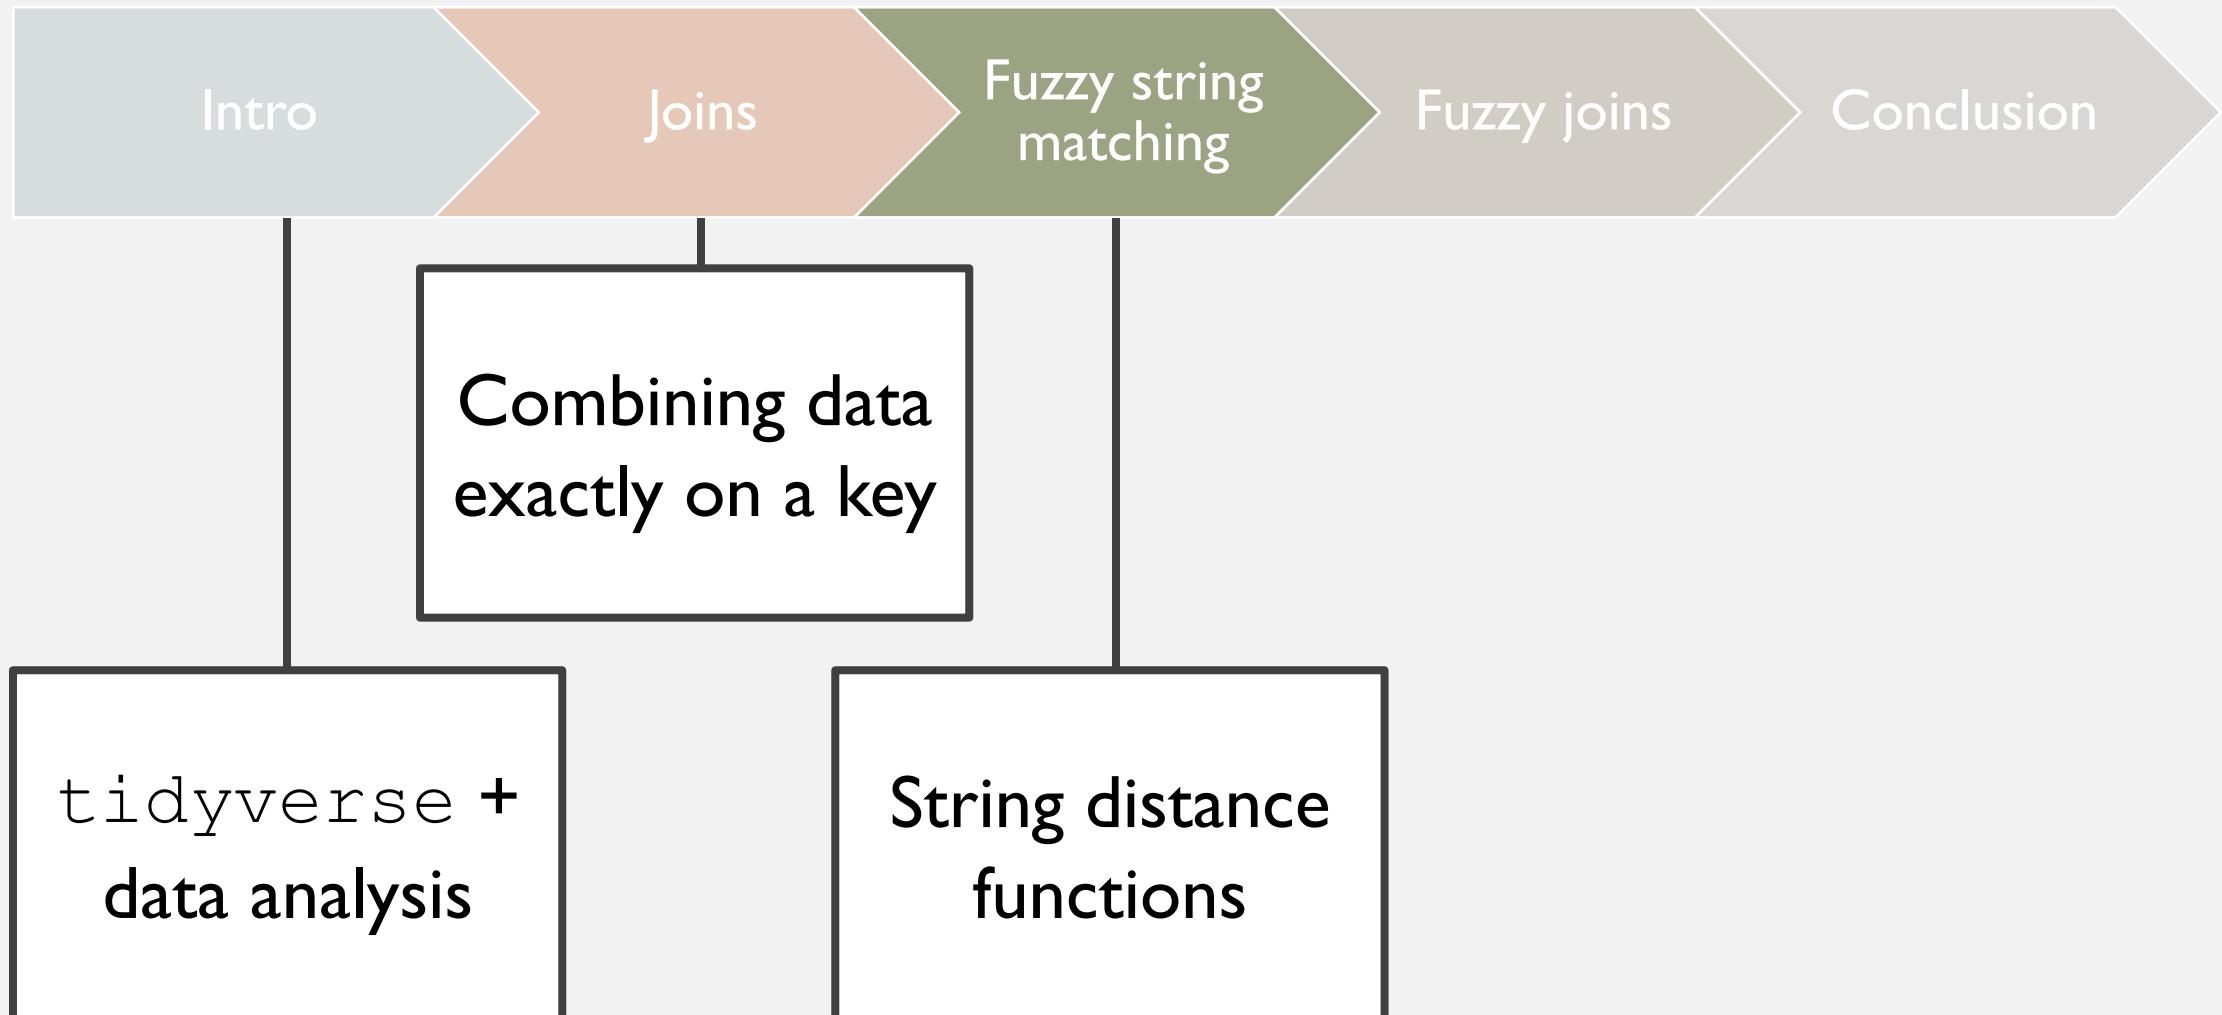


tidyverse +
data analysis

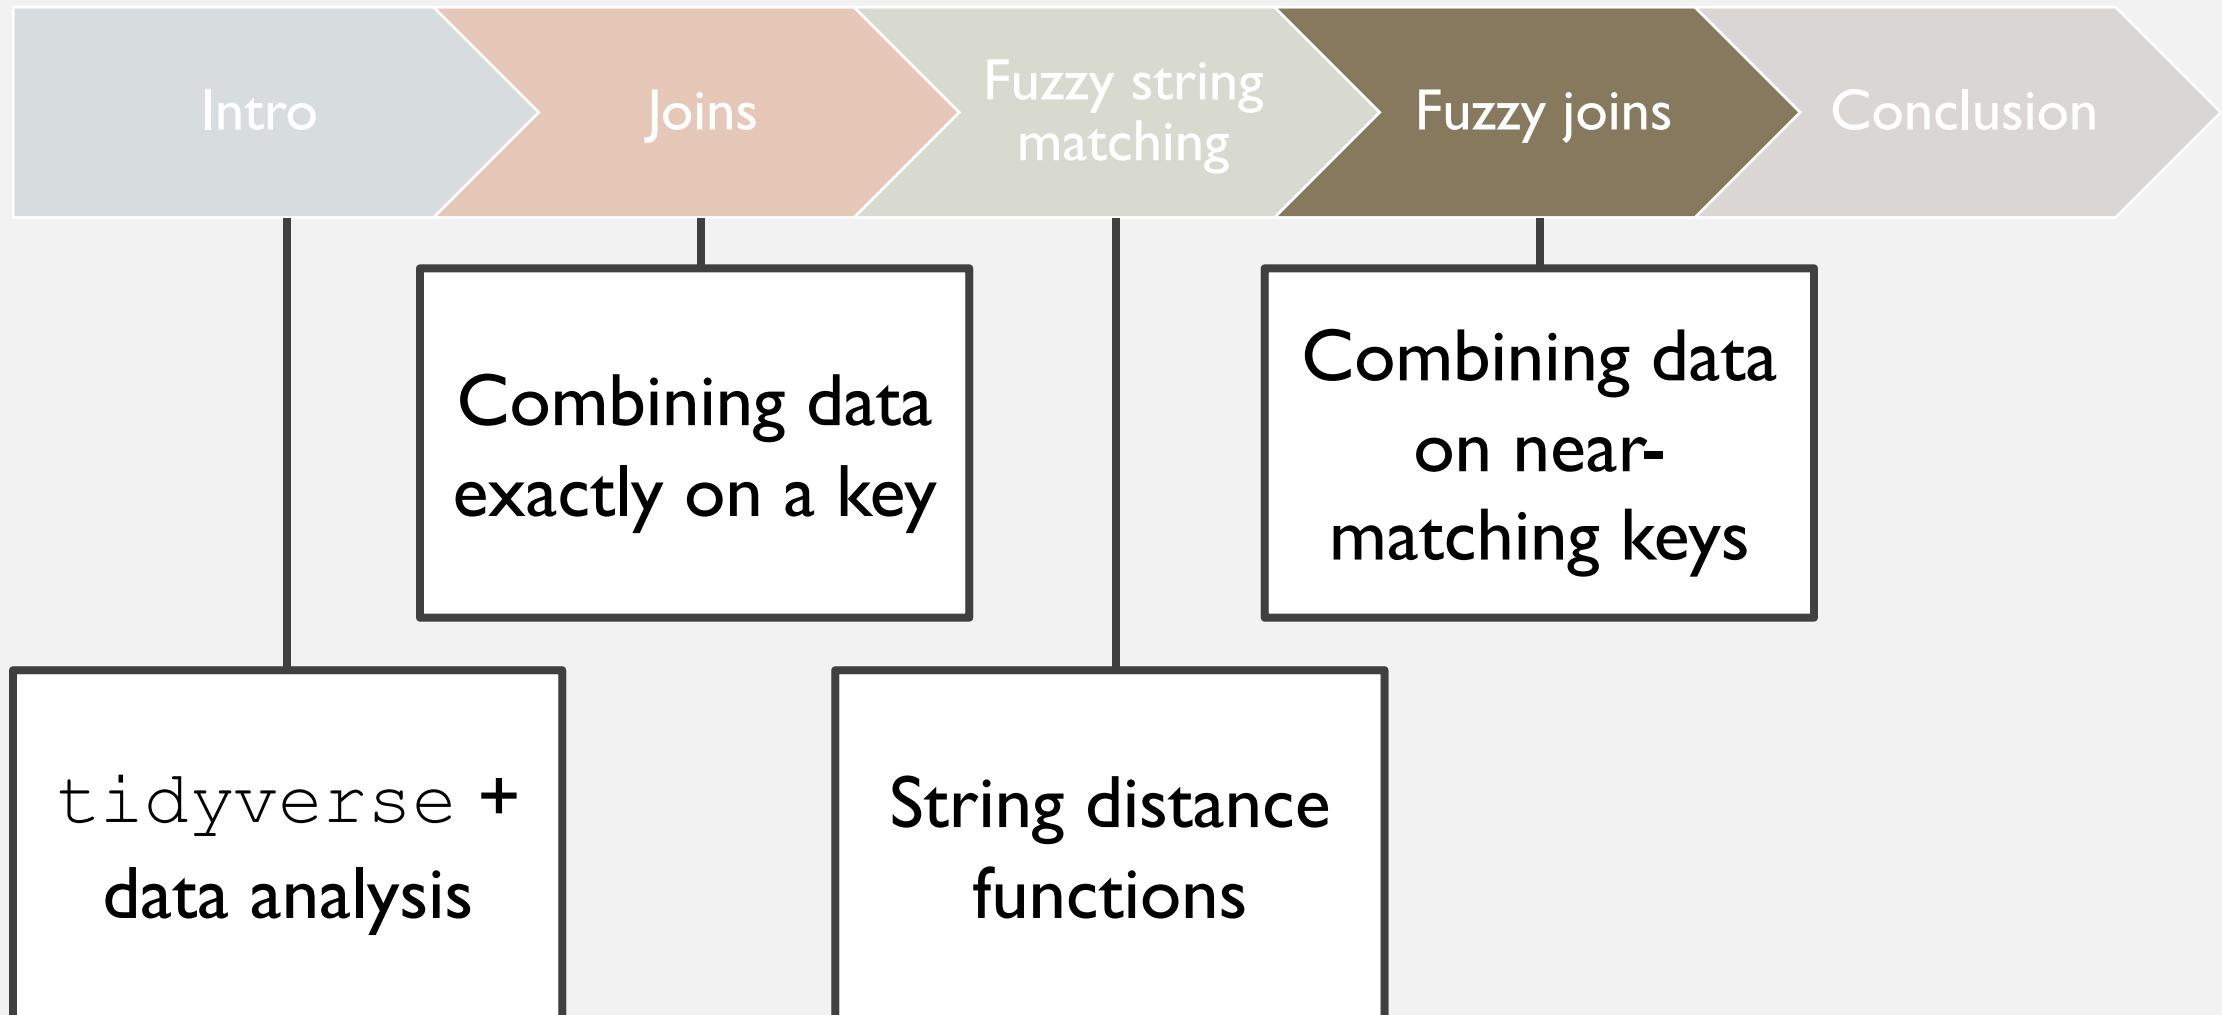
ORDER OF OPERATIONS



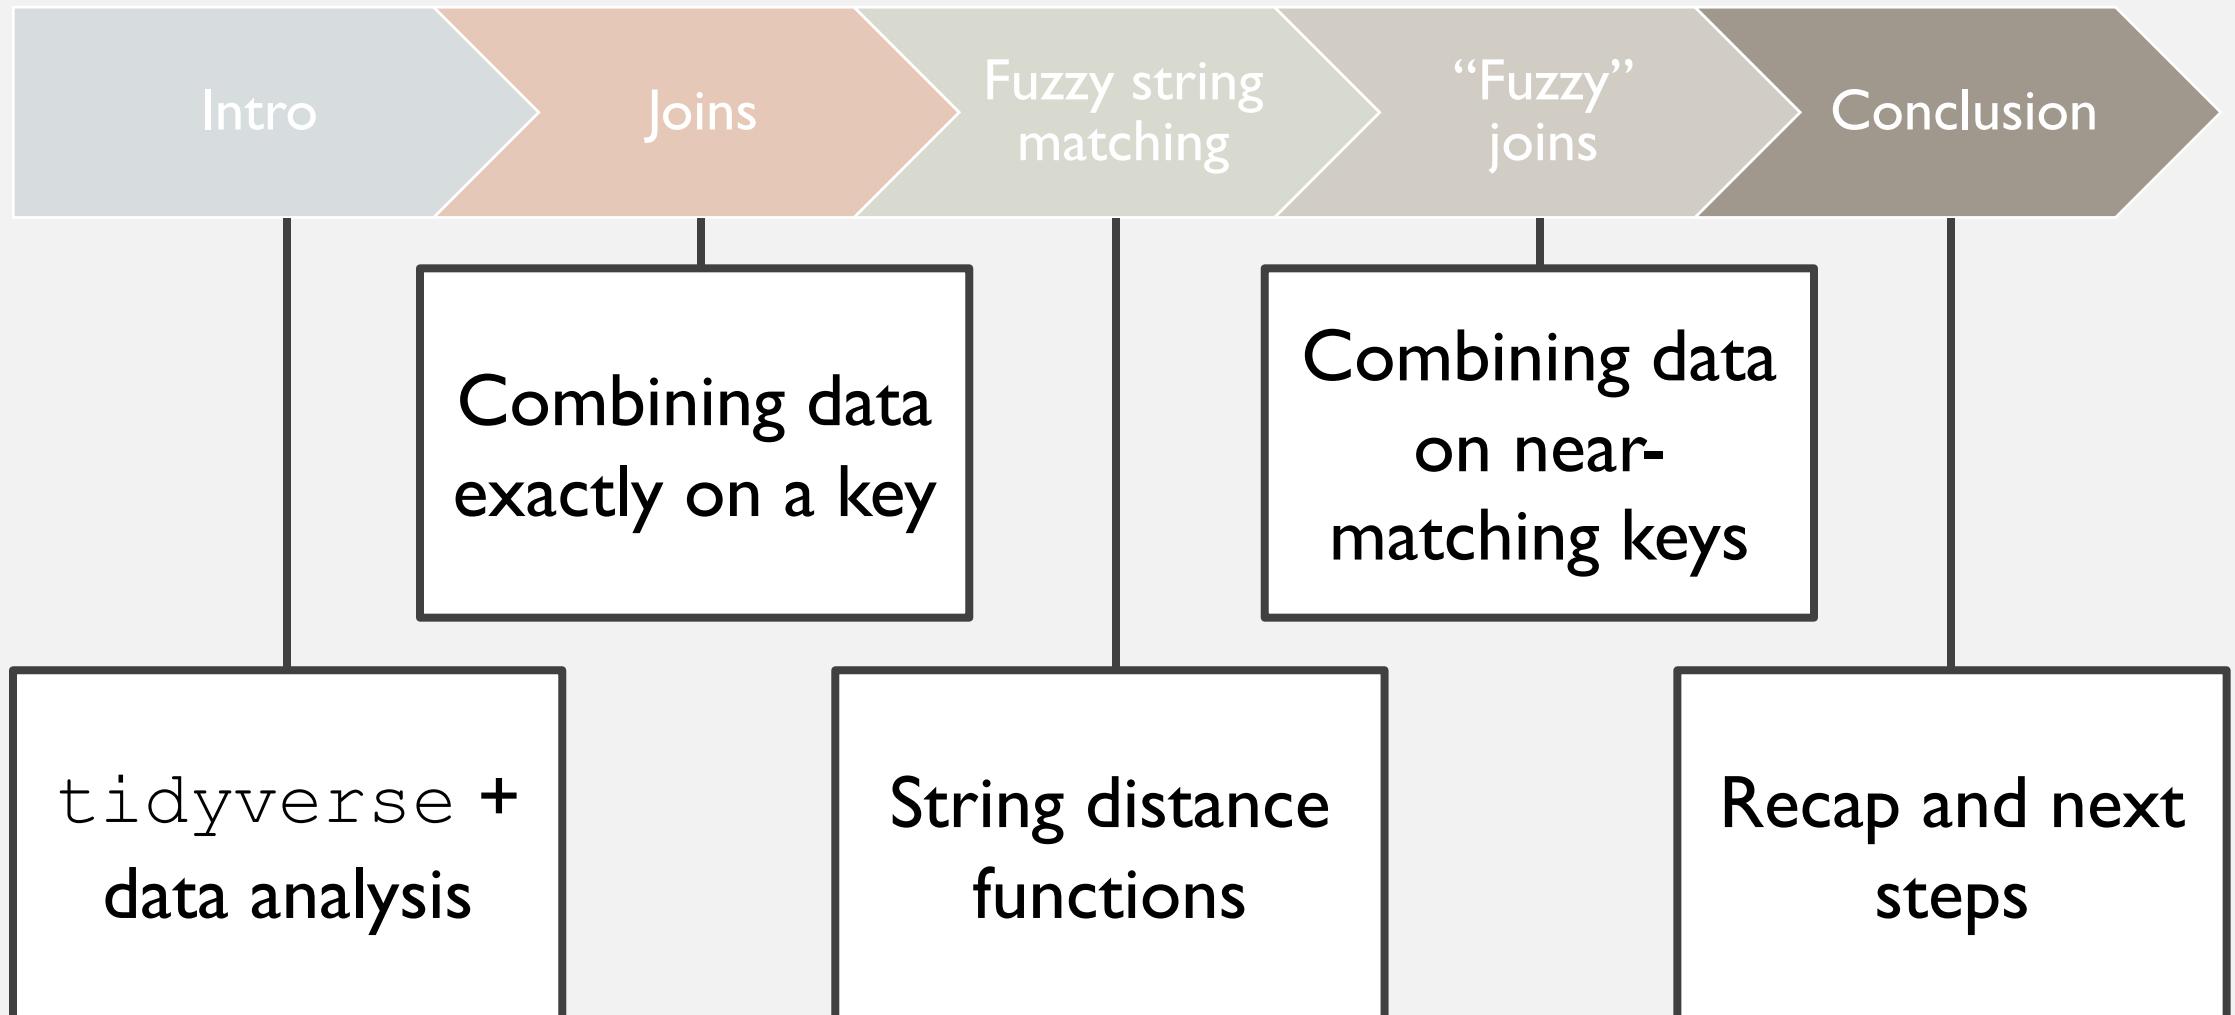
ORDER OF OPERATIONS



ORDER OF OPERATIONS



ORDER OF OPERATIONS



ANALYZING SCRAPED DATA

Intro

Joins

Fuzzy string
matching

Fuzzy joins

Conclusion

Today, we'll analyze a web-scraped manual entry data set. Some of the challenges associated with this kind of data include:

Today, we'll analyze a web-scraped manual entry data set. Some of the challenges associated with this kind of data include:

Spelling/keystroke
errors

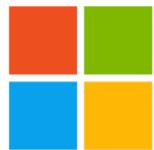


Gooooooooogle

Today, we'll analyze a web-scraped manual entry data set. Some of the challenges associated with this kind of data include:

Spelling/keystroke
errors

Multiple names
per entity



Microsoft



Apple

Gooooooooogle



Apple Inc.

Today, we'll analyze a web-scraped manual entry data set. Some of the challenges associated with this kind of data include:

Spelling/keystroke
errors



Multiple names/keys
per entity



Extra characters



Gooooooooogle



amazon



- Degree
- Subject
- Application year
- School name
- Decision
- Decision date
- Notes

| | | | | |
|------------------------------|---|--------------------------------------|--|--|
| Valdosta State University | Communication Sciences And Disorders, Masters (S20) | Accepted via Website on 15 Nov 2019 | 15 Nov 2019 | To the posted below - I am going in spring 2020 as well! I don't think there is a Facebook page for it yet |
| University Of Toronto | Geography, Masters (F19) | Rejected via E-mail on 2 Apr 2019 | I 14 Nov 2019 | :(sad |
| University Of Cambridge (UK) | Social Anthropology, PhD (F20) | Interview via E-mail on 28 Oct 2019 | 14 Nov 2019 | I was invited to interview by one of my potential advisors (I listed 2 in the application). Both of them joined in for a Skype interview since I am not in the UK (BA in the US, MA in Japan, currently in Japan). Lasted almost an hour but time went by quickly. Just waiting for the results now. |
| University Of Illinois | Bioinformatics, Masters (F20) | Other via Other on 15 Nov 2019 | I 14 Nov 2019 | Has Anyone received acceptance from them ? |
| Perimeter Institute | Physics, Masters (F20) | Other via Other on 14 Nov 2019 | I 14 Nov 2019 | Same here. Haven't heard anything from PI yet. One more day to go and hope we'll have good news. |
| University Of Iowa | Counselor Education And Supervision, PhD (S19) | Accepted via E-mail on 14 Nov 2019 ♦ | I 14 Nov 2019 | report spam |
| Cornell University | Engineering Management, Masters (S20) | Rejected via E-mail on 14 Nov 2019 ♦ | Undergrad GPA: 3.80 GRE General (V/Q/W): 157/170/3.50 GRE Subject: n/a | to the rest of you that applied for it :) |
| Georgia Tech | Operations Research, MC | Other via Other on 14 Nov 2019 | 14 Nov 2019 | Has any one heard back from GA Tech for the Online Operations Research Masters program? Submission deadline was 9/1/19 and my app still says "In Progress". |

```

# Load libraries and functions.
library(tidyverse)
library(stringdist)
library(fuzzyjoin)
source("scripts/user_defined_functions.R")

##### 1: ANALYZING SCRAPED DATA
#####
#####

# Table 2: admission results table
admissions_results <- read_csv("data/gradcafe_cs_results.csv")

admissions_results

```

```

> admissions_results
# A tibble: 42,613 x 11
  institution      date_added decision decision_date decision_day program degree gpa gre_verbal gre_quant gre_awa
  <chr>           <chr>       <chr>     <date>        <chr>      <chr>   <chr> <dbl>    <dbl>      <dbl>      <dbl>
1 University Of Sou... 11 Nov 20... Accepted 2019-11-11 Monday Computer... Maste... NA          NA        NA
2 University Of Mas... 11 Nov 20... Rejected 2019-11-11 Monday Computer... Maste... NA          NA        NA
3 UMass-Amherst      10 Nov 20... Rejected 2019-11-11 Monday Computer... Maste... NA          NA        NA
4 Columbia Universi... 10 Nov 20... Other    2019-11-11 Monday Computer... Maste... NA          NA        NA
5 University Of Mas... 10 Nov 20... Rejected 2019-11-09 Saturday Computer... Maste... NA          NA        NA

```

Exercise: use tidyverse-compliant functions to investigate the admissions results data. (Jump to script)

- Compute the mean GPA ($0 < \text{gpa} \leq 4$) for PhD applicants only
- Compute the mean *total* GRE scores ($130 \leq \text{gre_*} \leq 170$) for PhD applicants only
- Print the 50 schools with the most results overall in descending order

COMBINING DATA SETS

Intro

Joins

Fuzzy string
matching

Fuzzy joins

Conclusion

Best Computer Science Schools

Ranked in 2018, part of [Best Science Schools](#)

Earning a graduate degree in computer science can lead to positions in research institutions, government agencies, technology companies and colleges and universities. These are the top computer science schools. Each school's score reflects its average rating on a scale from 1 (marginal) to 5 (outstanding), based on a survey of academics at peer institutions. [Read the methodology »](#)



SUMMARY ▾



188 schools

[Computer Science Schools X](#)

[CLEAR ALL](#)

SORT BY: Rank (high to low) ▾

PROGRAM RANKINGS ^

Sciences ▾

Computer Science ▾

All Specialties ▾

SCHOOL NAME ^

School Name

LOCATION ^

| NAME/RANK | PEER ASSESSMENT SCORE |
|---|-----------------------|
| Carnegie Mellon University Pittsburgh, PA 💡 #1 in Computer Science (tie) | 5.0 |
| Massachusetts Institute of Technology Cambridge, MA 💡 #1 in Computer Science (tie) | 5.0 |
| Stanford University Stanford, CA | 5.0 |

```
# Table 1: rankings table
rankings <- read_csv("data/rankings_table.csv")
```

```
rankings
```

| | institution | usnwr_rank | csr_rank | nrc_rank |
|----|--|------------|----------|----------|
| | <chr> | <dbl> | <dbl> | <dbl> |
| 1 | Carnegie Mellon University | 1 | 1 | 4 |
| 2 | Massachusetts Institute of Technology | 2 | 2 | 2 |
| 3 | Stanford University | 3 | 4 | 1 |
| 4 | University of California, Berkeley | 4 | 5 | 5 |
| 5 | University of Illinois, Urbana-Champaign | 5 | 3 | 6 |
| 6 | Cornell University | 6 | 7 | 7 |
| 7 | University of Washington | 7 | 6 | 24 |
| 8 | Georgia Institute of Technology | 8 | 11 | 12 |
| 9 | Princeton University | 9 | 20 | 3 |
| 10 | University of Texas, Austin | 10 | 17 | 14 |

Goal: we want to combine the rankings from the first table with the admissions results from the second table where the institution name matches into a single table – otherwise known as a **join** or **inner join**.

```
inner_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"),  
...)
```

x: name of first table

y: name of second table

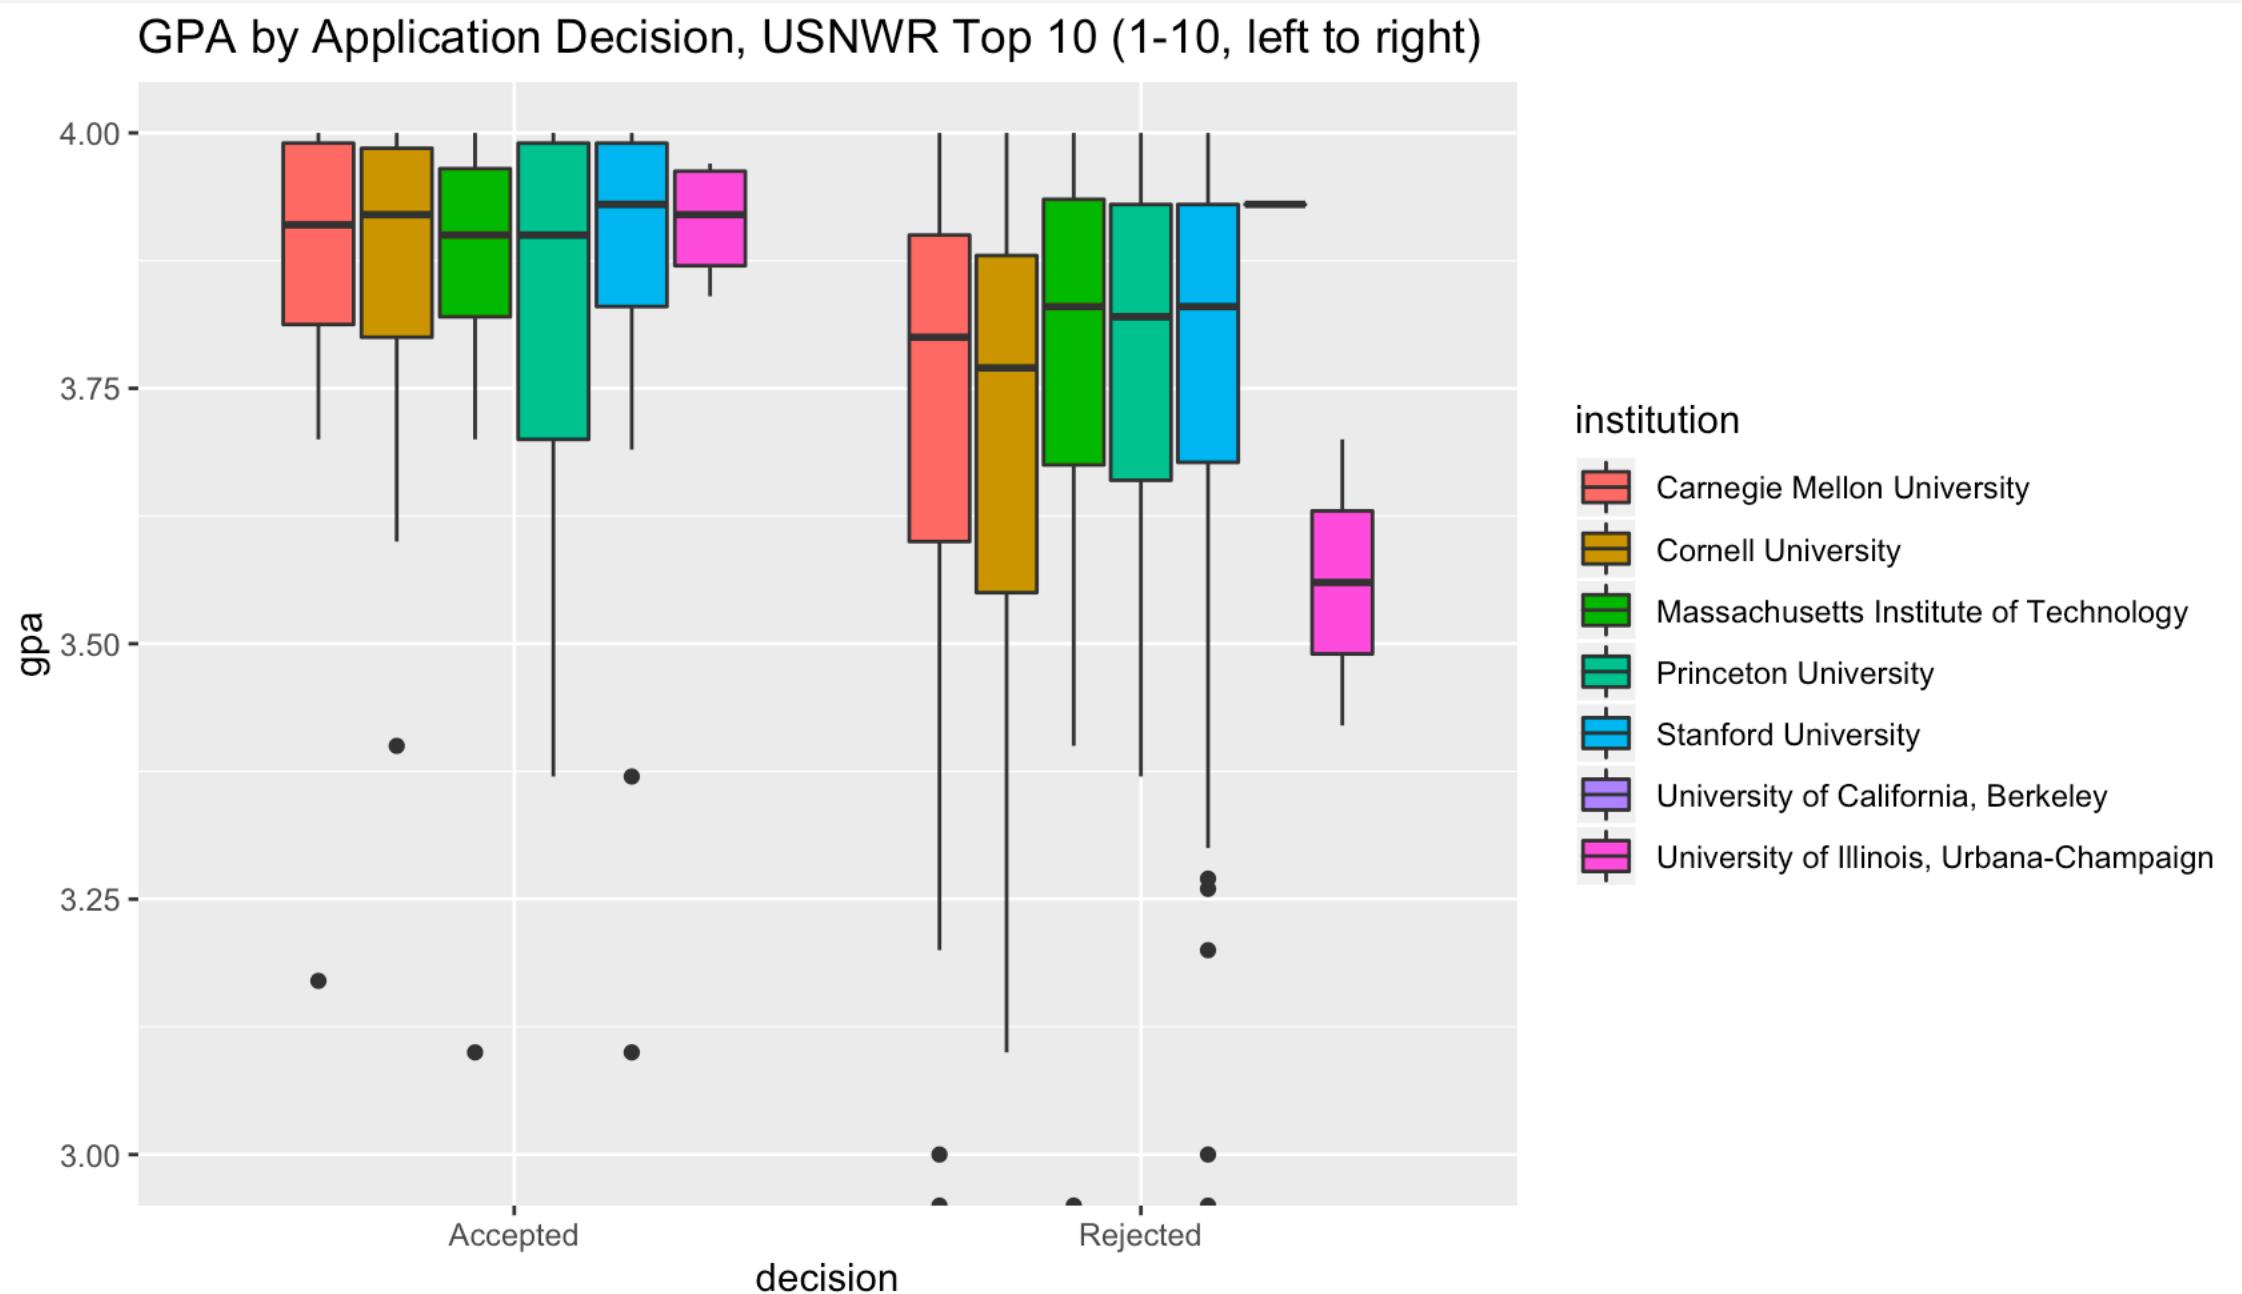
by: the name(s) of the **key(s)**, or column with values to be matched between x and y

suffix: suffix to distinguish non-key columns with names in both x and y

```
# INNER JOIN
join_results <- inner_join(rankings, admissions_results,
                           by = "institution")
join_results
```

```
> join_results <- inner_join(rankings, admissions_results,
+                             by = "institution")
> join_results
# A tibble: 3,268 x 14
  institution unswr_rank csr_rank nrc_rank date_added decision decision_date decision_day program degree   gpa
  <chr>           <dbl>    <dbl>    <dbl> <chr>      <chr>     <date>      <chr>    <chr> <chr>    <dbl>
1 Carnegie M...       1         1        4 3 May 2019 Rejected 2019-02-28 Thursday Comput... PhD     3.88
2 Carnegie M...       1         1        4 4 Apr 2019 Rejected 2019-04-04 Thursday Comput... Maste... NA
3 Carnegie M...       1         1        4 3 Apr 2019 Accepted 2019-04-03 Wednesday Comput... Maste... NA
4 Carnegie M...       1         1        4 3 Apr 2019 Accepted 2019-04-03 Wednesday Comput... Maste... 3.98
5 Carnegie M...       1         1        4 2 Apr 2019 Accepted 2019-04-02 Tuesday  Comput... Maste... 3.84
6 Carnegie M...       1         1        4 2 Apr 2019 Accepted 2019-04-02 Tuesday  Comput... Maste... NA
7 Carnegie M...       1         1        4 1 Apr 2019 Accepted 2019-04-01 Monday   Comput... Maste... NA
8 Carnegie M...       1         1        4 28 Mar 20... Other    2019-03-28 Thursday Comput... Maste... NA
9 Carnegie M...       1         1        4 27 Mar 20... Other    2019-03-28 Thursday Comput... Maste... NA
10 Carnegie M...      1         1        4 27 Mar 20... Other   2019-03-27 Wednesday Comput... Maste... NA
# ... with 3,258 more rows, and 3 more variables: gre_verbal <dbl>, gre_quant <dbl>, gre_awa <dbl>
```

Visualizing the join, we see only 7 institutions matched. What gives?



Let's look at keys containing "Berkeley" in admissions_results:

```
# Print all institutions in admission_results with substring "Berkeley"
admissions_results %>%
  filter(str_detect(institution, "Berkeley")) %>%
  distinct(institution) %>%
  print(n = nrow(.))
```

We find 38 different keys referencing “Berkeley”! (And these are just the ones that spelled “Berkeley” right *somewhere*.) Here are the first 12:

```
# A tibble: 38 x 1
  institution
  <chr>
  1 University Of California, Berkeley
  2 UC Berkeley
  3 Berkeley
  4 University Of California, Berkeleyley
  5 University Of California Berkeley
  6 Univeristy Of California, Berkeley
  7 UC Berkeley (UCB)
  8 University Of California Berkeley (UCB)
  9 University Of California Berkeley ( UCB )
 10 University Of California Berkeley (UC Berkeley)
 11 Computer Science Berkeley
 12 University Of California, Berkeley (UCB)
```

Standardization: we can fix many of the problems with the keys by...

- conversion to lower case
- stripping punctuation
- cleanup of leading/trailing/multiple whitespaces
- removing relatively uninformative words (e.g. words like **university, institute, school**; also **at** and **of**)
 - why? common/structural words artificially reduce distance of mismatches and offer little extra info for true matches

A user-defined function loaded by `source()` does all of this for us.

Adding standardized keys and redoing the join (jump to script):

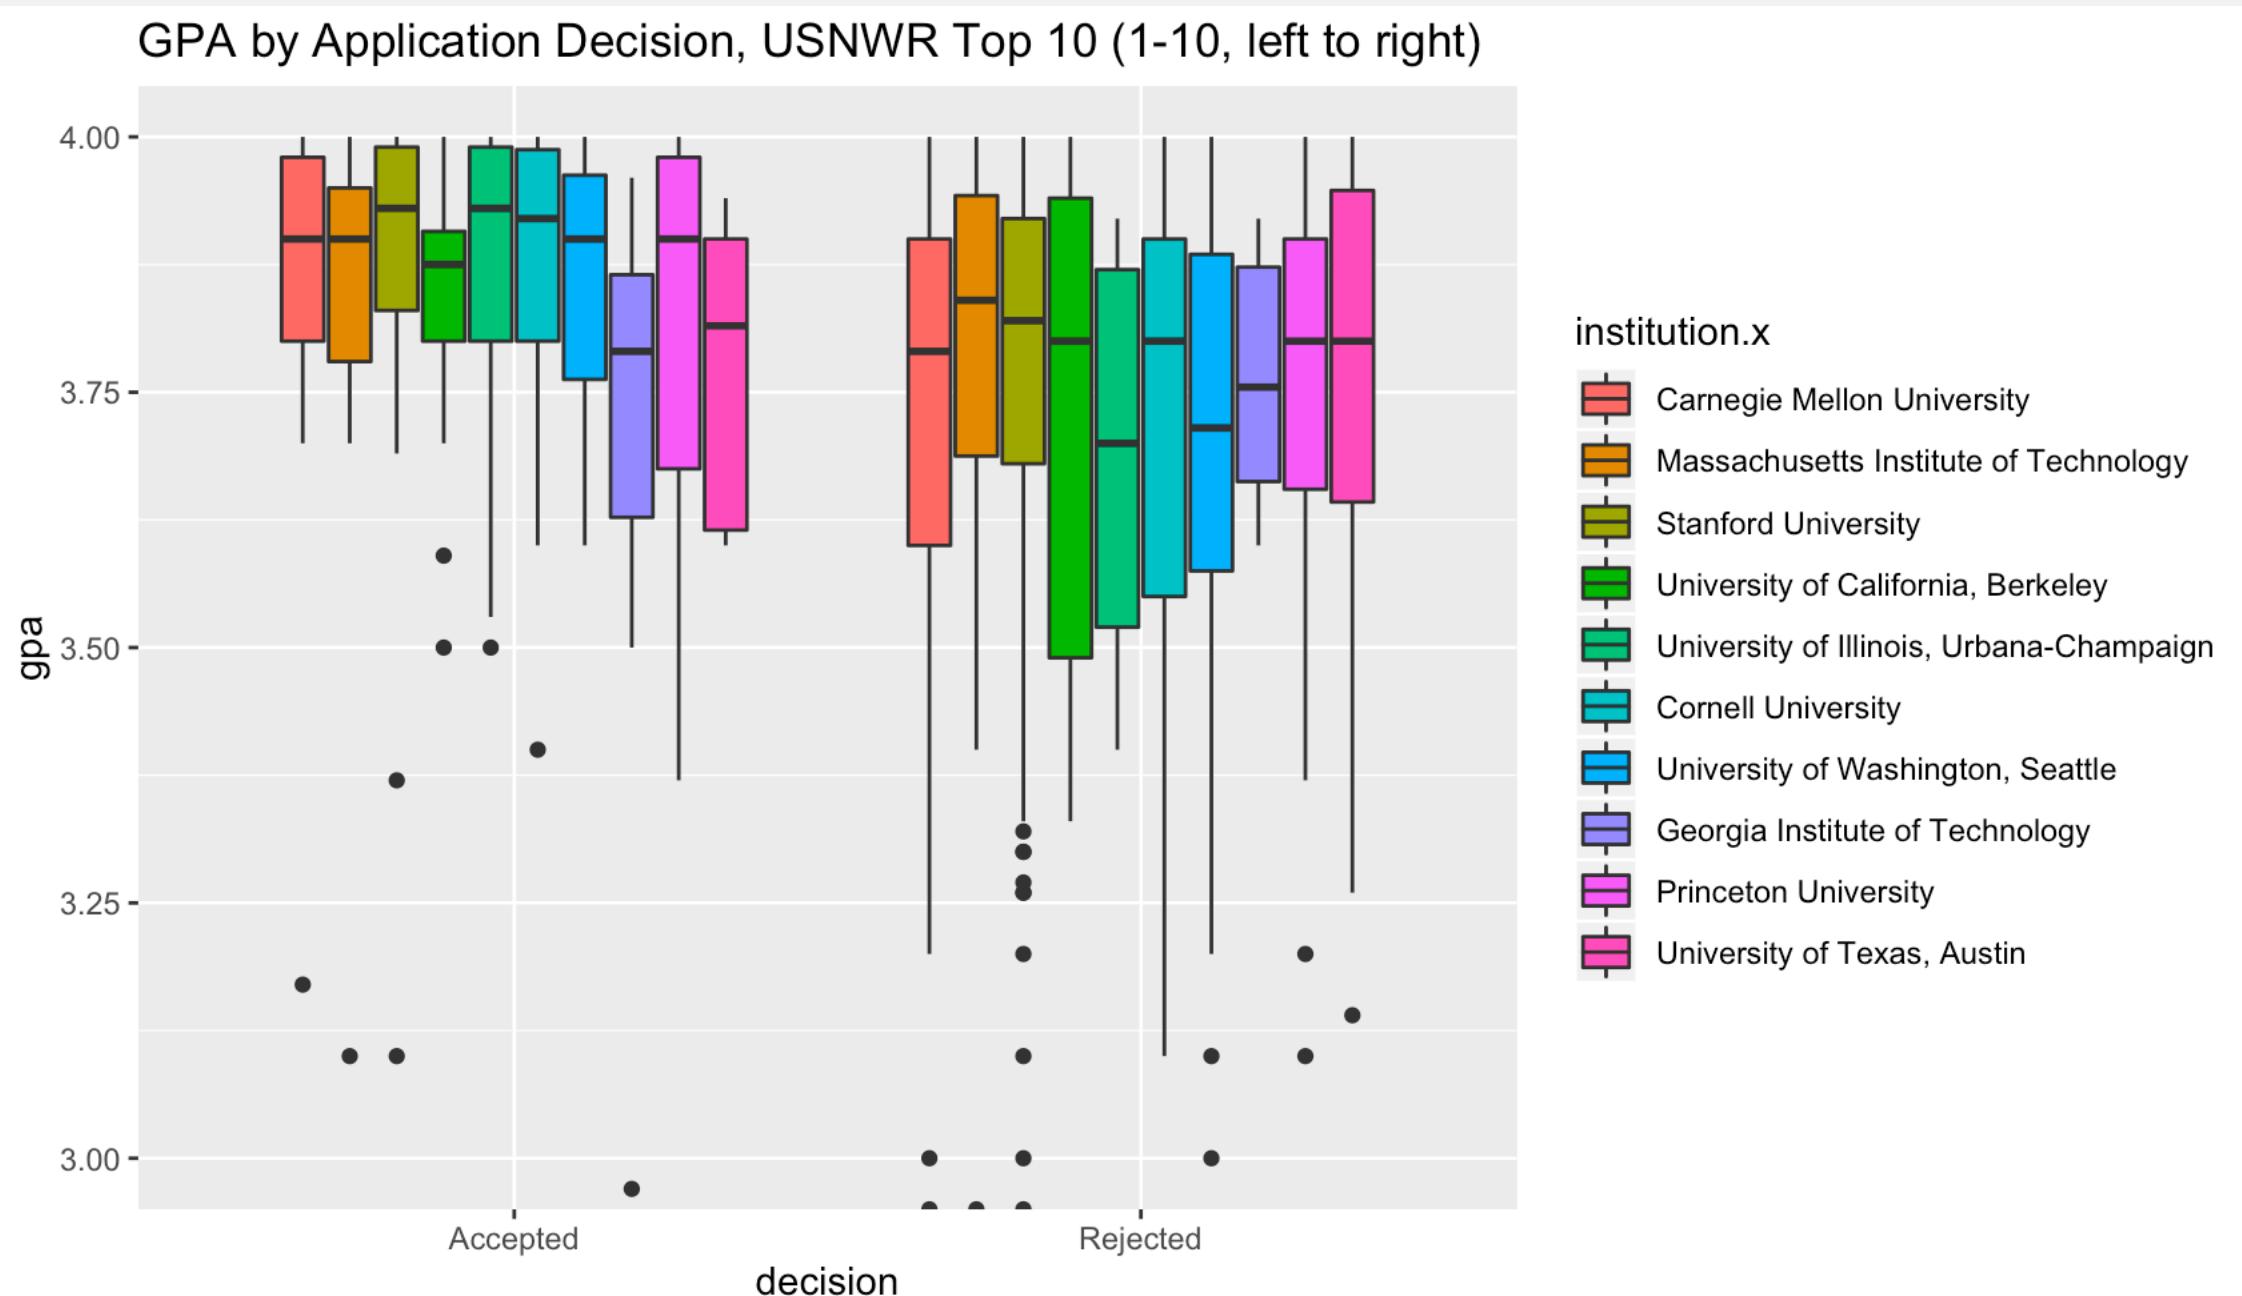
```
# ADD A STANDARDIZED KEY TO BOTH TABLES
rankings_adj <- rankings %>%
  mutate(institution_key = standardize_key(institution))

admissions_results_adj <- admissions_results %>%
  mutate(institution_key = standardize_key(institution))

# INNER JOIN ON STANDARDIZED KEY
join_results_adj <- inner_join(rankings_adj, admissions_results_adj,
                                 by = "institution_key")
join_results_adj
```

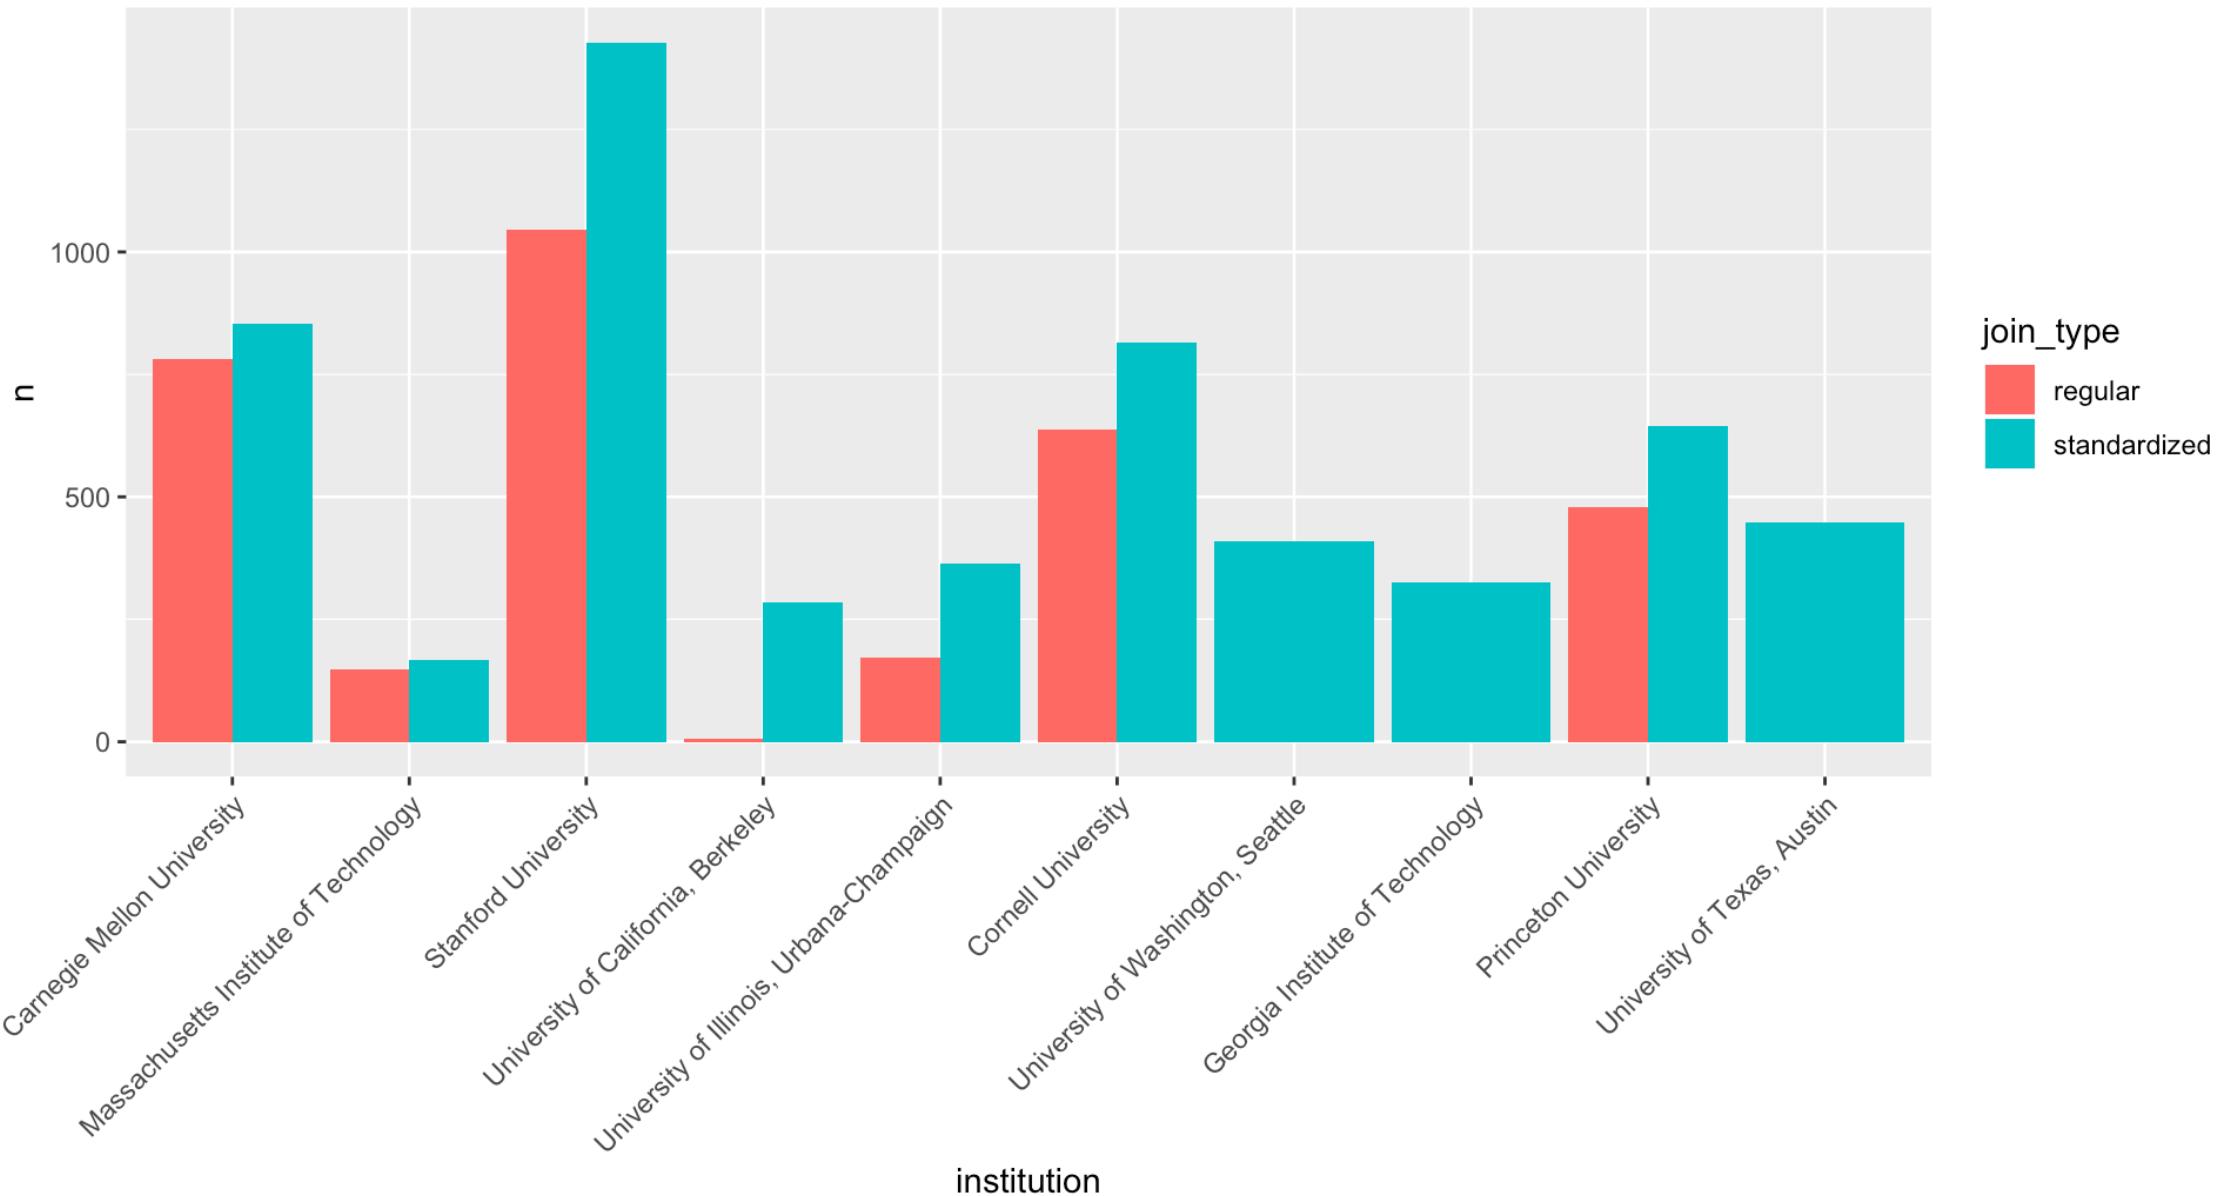
How many more matches do we see?

This looks better – at least all 10 institutions match some results.



Idiosyncratic number of matches suggests more room for improvement.

Number of matches by join type and institution



MEASURING CLOSENESS OF STRINGS

Intro

Joins

Fuzzy string
matching

Fuzzy joins

Conclusion

The difference between two strings can be measured by a string distance function. These functions come in many different flavors and are useful in different contexts – we'll discuss four of them.

```
stringdist(a, b, method = c("osa", "lv", "dl", "hamming", "lcs", "qgram",
  "cosine", "jaccard", "jw", "soundex"), useBytes = FALSE, weight = c(d
= 1, i = 1, s = 1, t = 1), q = 1, p = 0, bt = 0,
nthread = getOption("sd_num_thread"))
```

a = first string

b = second string

method = distance function name

weight, q, p, bt = auxiliary arguments for methods

I. Levenshtein distance: given two strings, at least how many character insertions, deletions, and substitutions would it take to transform the first string to a second string?

Start with “Northwestern”.

“Northwestern”

distance = 0

I. Levenshtein distance: given two strings, at least how many character insertions, deletions, and substitutions would it take to transform the first string to a second string?

Start with “Northwestern”. Substitute **w** for **e**.

“North**w**estern” → “North**e**stern”

distance = 1

I. Levenshtein distance: given two strings, at least how many character insertions, deletions, and substitutions would it take to transform the first string to a second string?

Start with “Northwestern”. Substitute **w** for **e**. Substitute **e** for **a**.

“Northwestern” → “Northe**e**stern” → “Northe**a**stern”

distance = 2

Check for yourself:

```
stringdist("Northwestern", "Northeastern", method = "lv")
```

2. Levenshtein distance ratio: given two strings, compute the Levenshtein distance. Then, divide by the length of the longer string.

“Northwestern” → “Northeestern” → “Northeastern”

2. Levenshtein distance ratio: given two strings, compute the Levenshtein distance. Then, divide by the length of the longer string.

“Northwestern” → “Northeastern” → “Northeastern” = 2

`str_length(“Northwestern”)` = `str_length(“Northeastern”)` = 12

2. Levenshtein distance ratio: given two strings, compute the Levenshtein distance. Then, divide by the length of the longer string.

“Northwestern” → “Northeastern” → “Northeastern” = 2

`str_length("Northwestern") = str_length("Northeastern") = 12`

so, the distance ratio is $2/12 = .167$

Check for yourself:

```
lvr("Northwestern", "Northeastern")
```

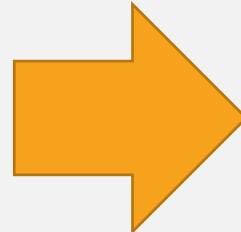
3. q-gram distance: given two strings, divide them into *grams* of size q , throwing out duplicates. How many grams are *not* shared by the strings?

($q = 2$)

Northwestern

2-grams

No



3. q-gram distance: given two strings, divide them into *grams* of size q , throwing out duplicates. How many grams are *not* shared by the strings?

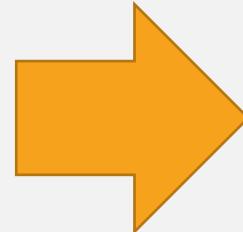
($q = 2$)

Northwestern

Northwestern

2-grams

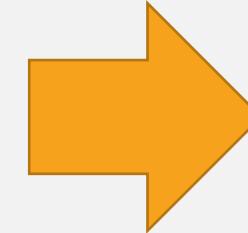
No or



3. q-gram distance: given two strings, divide them into *grams* of size q , throwing out duplicates. How many grams are *not* shared by the strings?

($q = 2$)

Northwestern



2-grams

No st

or hw

rt we

th es

te

rn

er

3. q-gram distance: given two strings, divide them into grams of size q , throwing out duplicates. How many grams are *not* shared by the strings?

| Northwestern | | | | Northeastern | | | |
|--------------|----|----|----|--------------|----|----|----|
| No | or | rt | th | No | or | rt | th |
| | hw | | es | | he | | as |
| st | | we | | st | | ea | |
| te | | | rn | te | | | rn |
| er | | | | er | | | |

distance = 6

Check for yourself:

```
stringdist("Northwestern", "Northeastern", method = "qgram", q = 2)
```

4. Jaccard distance: given two strings, compute the q -gram distance, and divide by the number of unique grams between the two strings.

| | | |
|----|----|----|
| es | hw | ea |
| we | he | as |

4. Jaccard distance: given two strings, compute the q -gram distance, and divide by the number of unique grams between the two strings.

| | | | | |
|-------|----|----|----|----|
| | es | hw | ea | |
| | we | he | as | |
| <hr/> | | | | |
| No | or | th | he | as |
| | hw | rt | es | |
| st | | we | ea | |
| te | er | rn | | |

4. Jaccard distance: given two strings, compute the q -gram distance, and divide by the number of unique grams between the two strings.

$$\frac{\text{es} \quad \text{hw} \quad \text{ea}}{\text{we} \quad \text{he} \quad \text{as}} = \frac{6}{14} = .429$$

No or th
st hw rt es he
te we rn ea

Check for yourself:

```
stringdist("Northwestern", "Northeastern", method = "jaccard", q = 2)
```

Notes:

The value of q should generally get larger as the keys get longer; $q = 2$ or 3 suffices for most strings of sentence length or shorter.

Increasing q generally means

- distances overall will increase – bigger library of grams, more mismatches
- you care more about the *structure* of the word than letters

If one string pair has a lower distance at $q = 2$ than another, that does not guarantee that pair will have a lower distance at $q = 3$

Exercise: find the closest match (smallest distance) among the pairs of strings below by (a) Levenshtein distance ratio; (b) Jaccard distance, $q = 2$.

Dartmouth College and **Darmtouth College**

University of California and **The University of California**

Stanford University and **University of Stanford**

How do the distances vary between the two metrics?

Given that these strings are all matches, do you prefer one metric?

COMBINING DATA SETS, FUZZILY

Intro

Joins

Fuzzy string
matching

Fuzzy joins

Conclusion

```
stringdist_join(x, y, by = NULL, max_dist = 2, method = c("osa",
  "lv", "dl", "hamming", "lcs", "qgram", "cosine", "jaccard", "jw",
  "soundex"), mode = "inner", ignore_case = FALSE,
  distance_col = NULL, ...)
```

x: name of first table

y: name of first table

by: the name(s) of the **key(s)**, or column with values to be fuzzily matched between x and y

max_dist: returns only matches with distance smaller than this threshold

method: specifies the string distance function to use

distance_col: the name of a column to be added with the key distance for each row

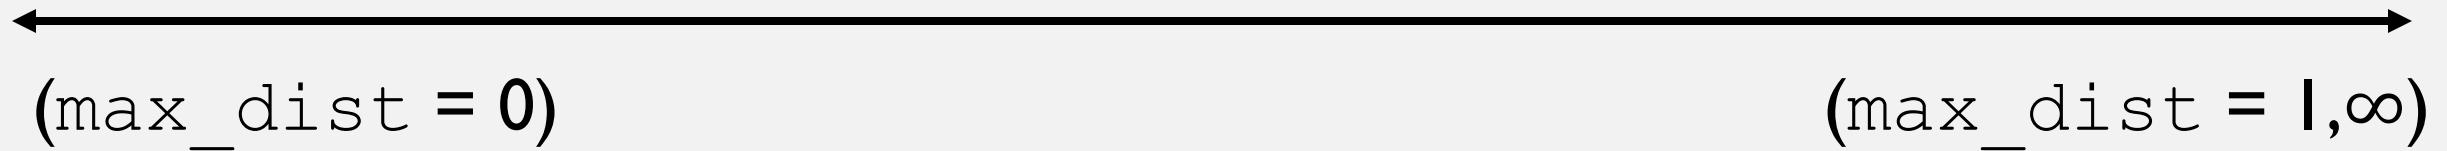
(Jump to R script)

```
##### 4: COMBINING DATA SETS FUZZILY
#####
#####

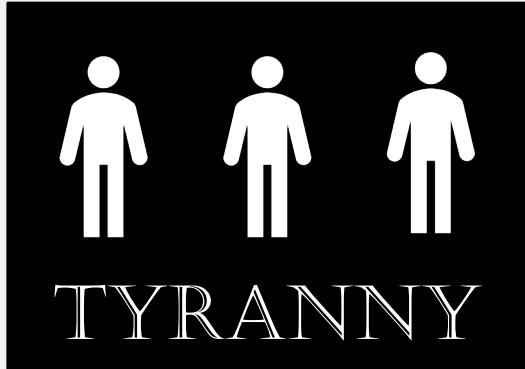
# run fuzzy join
fuzzy_join_results <- stringdist_inner_join(rankings_adj,
                                              admissions_results_adj,
                                              by = c("institution_key" = "institution_key"),
                                              method = "jaccard", q = 2, max_dist = .2,
                                              distance_col = "dist")
fuzzy_join_results

# see most different matches first
fuzzy_join_results %>%
  filter(dist > 0) %>%
  group_by(institution.x, institution.y) %>%
  summarize(num_results = n(),
            dist = first(dist)) %>%
  arrange(desc(dist)) %>%
  View()
```

Selecting a `max_dist`: assessing your appetite for potential mismatches.



Selecting a max_dist: assessing your appetite for potential mismatches.



↔ (max_dist = 0) → (max_dist = 1,∞)

- Only perfect matches
- No mismatches
(0 *false positives*, perfect *precision*)

Selecting a max_dist: assessing your appetite for potential mismatches.



TYRANNY



ANARCHY



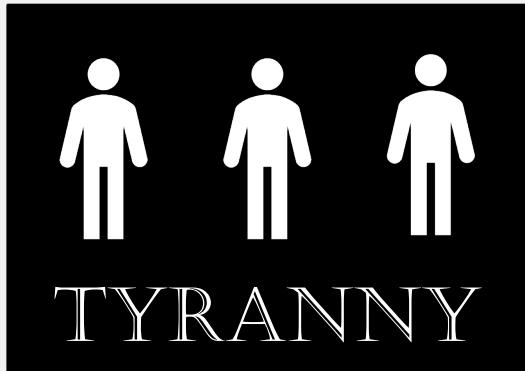
(max_dist = 0)

- Only perfect matches
- No mismatches
(0 *false positives*, perfect precision)

(max_dist = 1,∞)

- **EVERYTHING** matches!
- Lots of *false positives*, perfect recall

Selecting a max_dist: assessing your appetite for potential mismatches.



↔ (max_dist = 0)

- Only perfect matches
- No mismatches (*0 false positives, perfect precision*)

- Happy medium!
- Lots of correct matches, few mismatches
- Varies by **user** and **use case**

↔ (max_dist = 1,∞)

- **EVERYTHING** matches!
- Lots of *false positives, perfect recall*

A general strategy for finding your happy medium:

1. Start with near-tyranny (**low** `max_dist`, say .05)
2. Pick a *stopping rule*
3. Increase `max_dist` a bit, and look at the matches with the highest distance between them. Do they break your stopping rule?

Sample stopping rules:

- *I won't accept any mismatches at all (no risk, lowest reward)*
- *I'll accept at most X mismatches total (low risk, moderate reward)*
- *increasing max_dist must result in Y times as many true matches than mismatches (medium risk, highest reward)*

Setting the threshold at .2:

| | institution.x | institution.y | num_results | dist |
|----|--|--|-------------|------------|
| 1 | Carnegie Mellon University | Canegie Mellon | 1 | 0.20000000 |
| 2 | Carnegie Mellon University | Carnege Mellon | 1 | 0.20000000 |
| 3 | Carnegie Mellon University | Carnege Mellon University | 24 | 0.20000000 |
| 4 | Carnegie Mellon University | Carnegi Mellon University | 3 | 0.20000000 |
| 5 | University of Illinois, Urbana-Champaign | Illinois At Urbana-Champaign (UIUC)inois At Urbana-... | 5 | 0.20000000 |
| 6 | University of Illinois, Urbana-Champaign | Illinois At Urbana-Champaign (UIUC)s | 1 | 0.20000000 |
| 7 | Princeton University | Princeton U | 2 | 0.20000000 |
| 8 | Massachusetts Institute of Technology | Massachussets Institute Of Technology (MIT) | 6 | 0.19230769 |
| 9 | Georgia Institute of Technology | GaTech (Georgia Institute Of Technology) | 220 | 0.19047619 |
| 10 | Georgia Institute of Technology | Georgia Institute Of Technology – Georgia Tech – GT | 3 | 0.19047619 |
| 11 | Georgia Institute of Technology | Georgia Institute Of Technology (GaTech) | 15 | 0.19047619 |
| 12 | Georgia Institute of Technology | Georgia Institute Of Technology (GaTech | 7 | 0.19047619 |
| 13 | Georgia Institute of Technology | Georgia Institute Of Technology (gatech) | 1 | 0.19047619 |
| 14 | Georgia Institute of Technology | Georgia Institute Of Technology (Gatech) | 7 | 0.19047619 |
| 15 | Georgia Institute of Technology | Georgia Institute Of Technology (GaTech) | 163 | 0.19047619 |
| 16 | Georgia Institute of Technology | Georgia Institute Of Technology (GATech) | 4 | 0.19047619 |
| 17 | Georgia Institute of Technology | Georgia Institute Of Technology(GaTech) | 1 | 0.19047619 |
| 18 | Carnegie Mellon University | Carneigie Mellon University | 1 | 0.18750000 |

Exercise: find the threshold you are most comfortable with in our fuzzy join problem. You can choose to use one of the example stopping rules, or just come up with your own by feel.

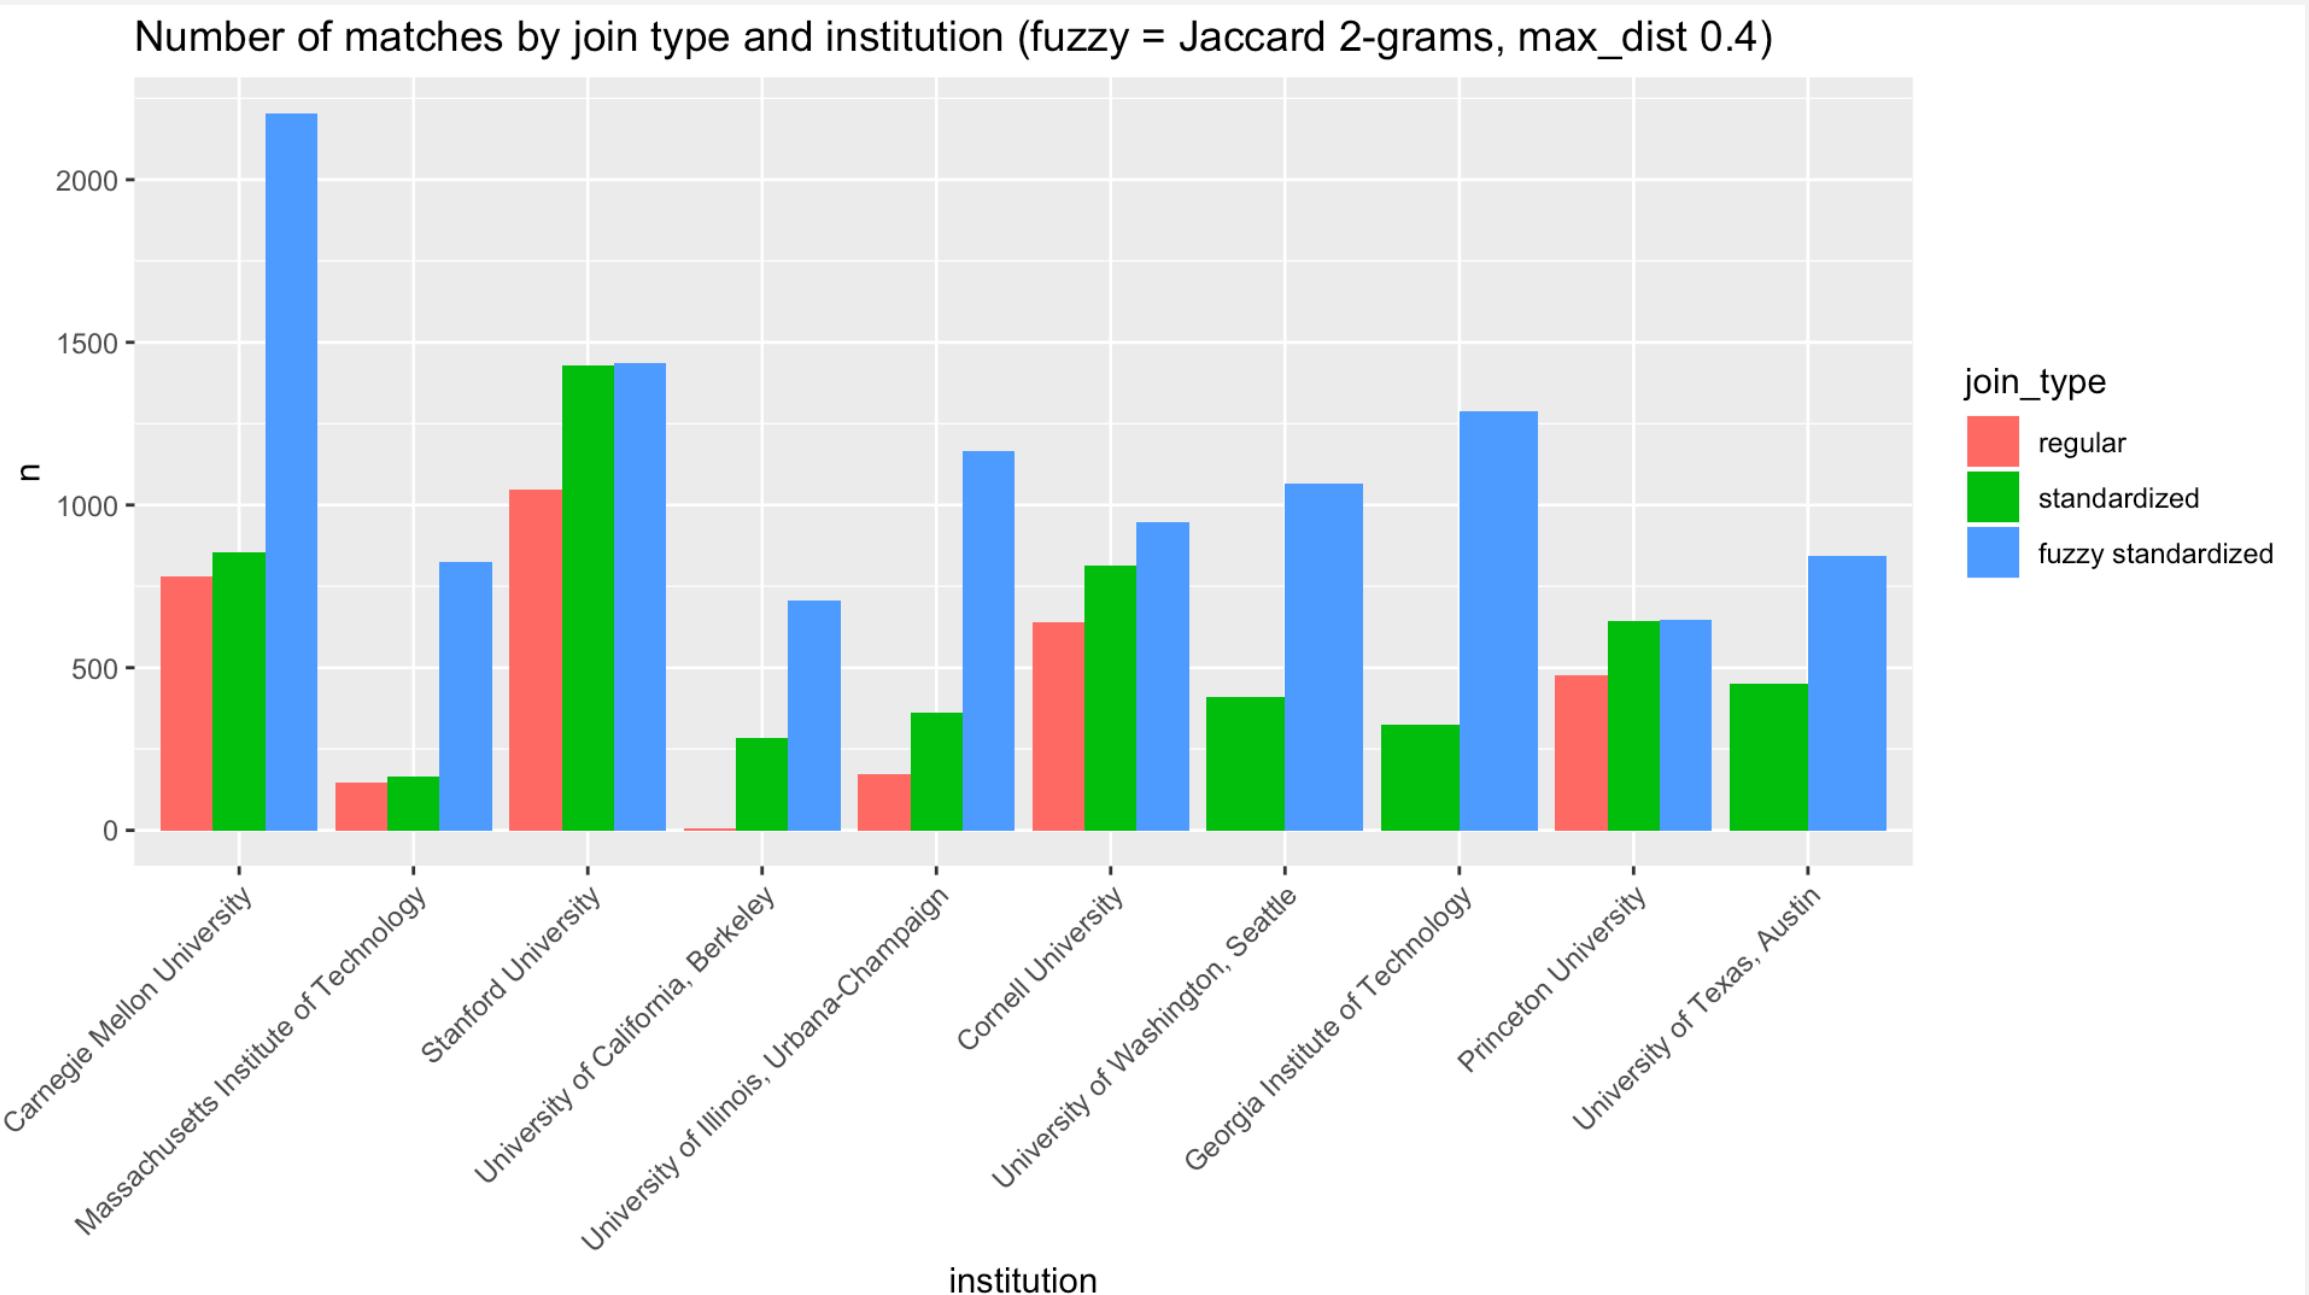
1. Rerun the join with a higher `max_dist` parameter
2. Check the matches with the largest distance using the data viewer
3. If you're still comfortable, repeat steps above and reassess

If you like: use whichever string distance you prefer. There's no built-in method for a join on Levenshtein ratio; I've written a wrapper that will do the join for you (`lvr_inner_join`), already loaded into your R session.

Setting the threshold at .518:

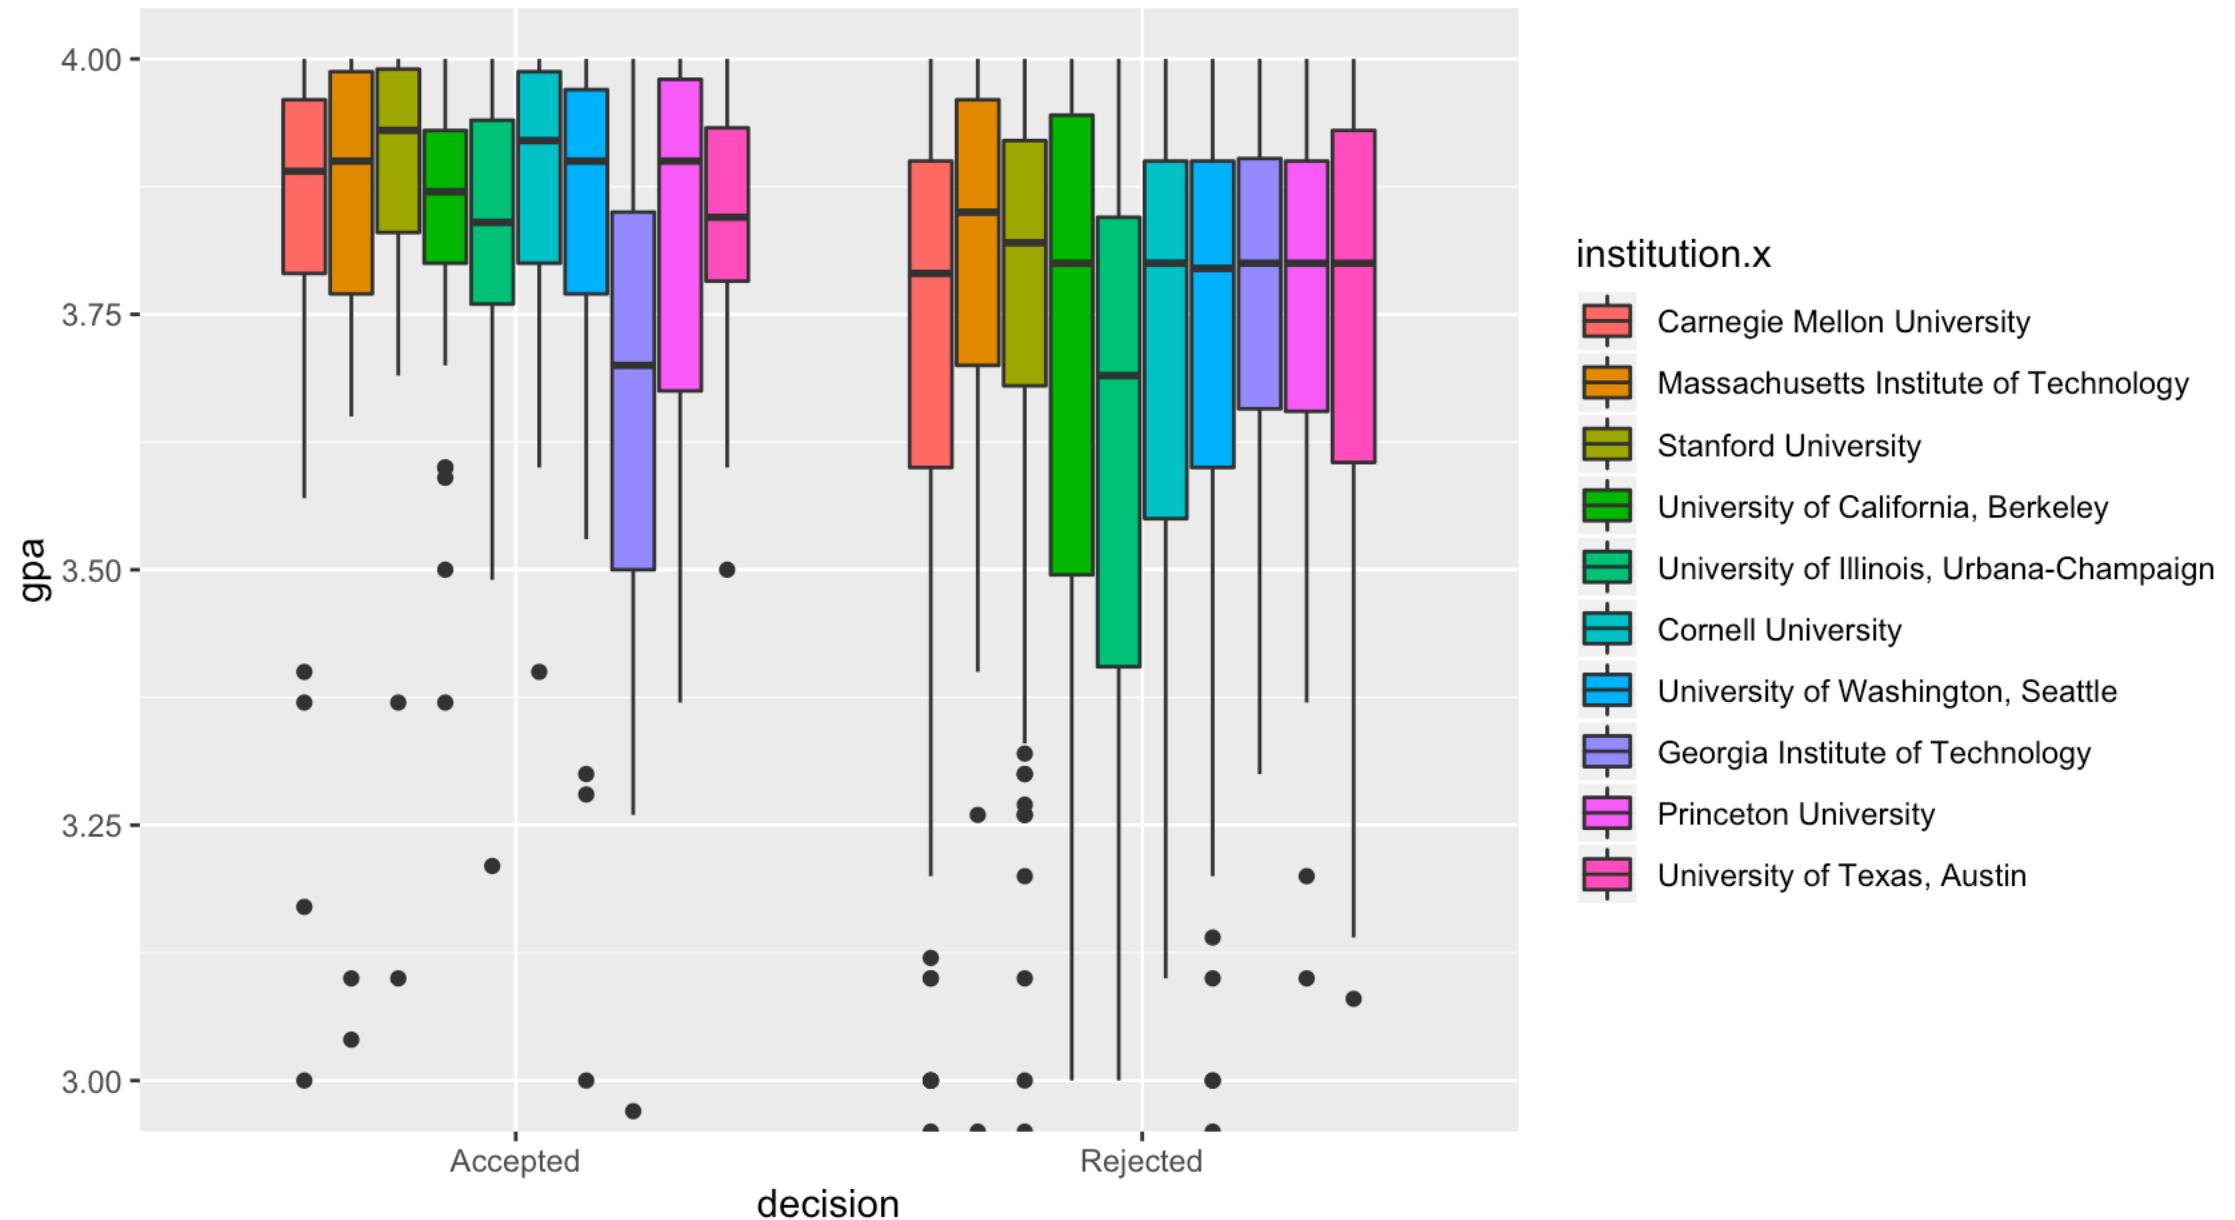
| | | | | |
|-----|---------------------------------------|--|-----|-----------|
| 179 | Georgia Institute of Technology | California Institute Of Technology (Caltech) | 68 | 0.5185185 |
| 180 | Carnegie Mellon University | Carnegie Mellon Univeristy (CMU) | 2 | 0.5172414 |
| 181 | Carnegie Mellon University | Carnegie Mellon University (CMU, CMU School Of Music) | 1 | 0.5172414 |
| 182 | Carnegie Mellon University | Carnegie Mellon University (CMU) – Robotics Instituteu | 1 | 0.5172414 |
| 183 | Carnegie Mellon University | Carnagie Mellon University (CMU) LTI | 15 | 0.5000000 |
| 184 | Massachusetts Institute of Technology | University Of Massachusett | 1 | 0.5000000 |
| 185 | University of California, Berkeley | University Of California | 2 | 0.5000000 |
| 186 | University of Washington | University Of Washington Milwaukee | 1 | 0.5000000 |
| 187 | University of Washington | University Of Washingtonhington (Seattle) | 1 | 0.5000000 |
| 188 | University of Washington | Uversity Of Washington | 1 | 0.5000000 |
| 189 | Georgia Institute of Technology | GeorgiaTech | 1 | 0.5000000 |
| 190 | Georgia Institute of Technology | Gorgia Tech | 2 | 0.5000000 |
| 191 | Georgia Institute of Technology | Indian Institute Of Technology | 1 | 0.5000000 |
| 6 | University of California, Berkeley | UC Berkeley | 365 | 0.6000000 |

Now we have much more data at minimal cost, and improved parity:



The colorful fruits of our labor!

GPA by Application Decision, USNWR Top 10 (1-10, left to right)



RECAP

Intro

Joins

Fuzzy string
matching

Fuzzy joins

Conclusion

RECAP

Today, we...

- discussed techniques for increasing the number of matches in a join via string distances
- compared a few choices of string distance for this purpose
- tuned parameters to maximize the true match rate while minimizing the false match rate (maximize recall at minimal cost to precision)
- worked through a real example on scraped data

CHALLENGE EXERCISE

In the scripts folder, take `home_exercise.R` loads the CS rankings and admissions data we just worked with, as well as rankings and scraped admissions data for **statistics** programs.

The exercise is to...

- use `stringdist_join` to match the each of the top 10 statistics and CS programs with their admissions results
- using `summarize`, compute mean GPA for **accepted PhD applicants** in each discipline
- using `inner_join`, compare institutions with top 10 programs in **both** disciplines – among applicants reporting acceptance, which discipline has a higher average GPA?

THANKS FOR YOUR TIME! QUESTIONS?